

Adam Zweiger

adamz@mit.edu | adamzweiger.github.io | github.com/adamzweiger

EDUCATION

Massachusetts Institute of Technology

Fall 2023 – Spring 2026

BS in Computer Science and Engineering, GPA: 5.0/5.0

- Relevant Coursework: **Theory:** Advanced Algorithms (G), Information Theory (G); **AI:** Deep Learning (G), Natural Language Processing (G); **Engineering:** Accelerated Computing (G), Distributed Systems (G), Operating Systems, Computer Systems Engineering, Software Construction
- Teaching Experience: 6.390 (Intro to ML) Lab Assistant

Carlmont High School

Fall 2019 – Spring 2023

GPA – 4.0/4.0

- Additional Stanford Online High School Coursework in Math and Physics: Linear Algebra, Differential Equations, Modern Algebra, Real Analysis, Complex Analysis, Modern Physics
- Euler Circle Advanced Classes: Markov Chains, Complex Analysis, Generating Functions

EXPERIENCE

Member of Technical Staff Intern | Cognition AI

January 2026 – April 2026

- Worked on reinforcement learning for coding models.

Undergraduate Researcher | MIT CSAIL

October 2024 – May 2026

- Research in language models and ML systems under Prof. Yoon Kim, with additional work under the guidance of Prof. Pulkit Agrawal and Prof. Jacob Andreas.
- Worked on data-efficient adaptation and long-context methods for language models.

Software Engineering Intern | Amazon Web Services

Summer 2024

PUBLICATIONS

Fast KV Compaction via Attention Matching. ICML 2026. **Zweiger, A.**, Fu X., Guo, H., Kim, Y.

- Introduced a new approach for fast and high-quality context compaction in latent space via Attention Matching.

Self-Adapting Language Models. NeurIPS 2025. **Zweiger, A.**, Pari, J., Guo, H., Akyürek, E., Kim, Y., Agrawal, P.

- Introduced SEAL, a meta-learning framework that teaches language models to generate optimal synthetic finetuning data and self-edit directives. Covered by WIRED and MIT News.

The Surprising Effectiveness of Test-Time Training for Few-Shot Learning. ICML 2025. Akyürek, E., Damani M., **Zweiger, A.**, Qiu, L., Guo, H., Pari, J., Kim, Y., Andreas, J.

- Demonstrated that combining test-time training with in-context learning yields state-of-the-art few-shot learning performance, achieving 62% on the ARC-AGI benchmark.

PROJECTS

HackMIT Organizer | Full-Stack Development (TypeScript, React, Flask, Redis, AWS)

2024 – 2025

- Worked on Pigeon, a RAG-powered email response client, and Plume, an application and judging platform for hackathons.
- Created CTF-style admissions puzzles.

HONORS & AWARDS

Putnam Top 200 (2023), USAMO Honorable Mention (2023), US Physics Camp (2022), USACO Platinum (2021), Canada/USA Mathcamp Attendee (2021, 2022), Atlas Fellow (\$50k Scholarship, 2022), MIT web.lab Winner (2024), ICML Gold Reviewer (2026)